

Chapter Two: Users, Files, and the Manual

Objectives

Ideas and Skills

- The role and use of online documentation
- The Unix file interface: open, read, write, lseek, close
- Reading, creating, and writing files
- File descriptors
- Buffering: user level, and kernel level

Ideas and Skills

- Kernel mode, user mode, and the cost of system calls
- How Unix represents time, how to format Unix time
- Using the utmp file to find the list of current users
- Detecting and reporting errors in system calls

System Calls and Functions

- Open, read, write, creat, lseek, close
- perror

Commands

- man
- who
- cp
- login

The who command

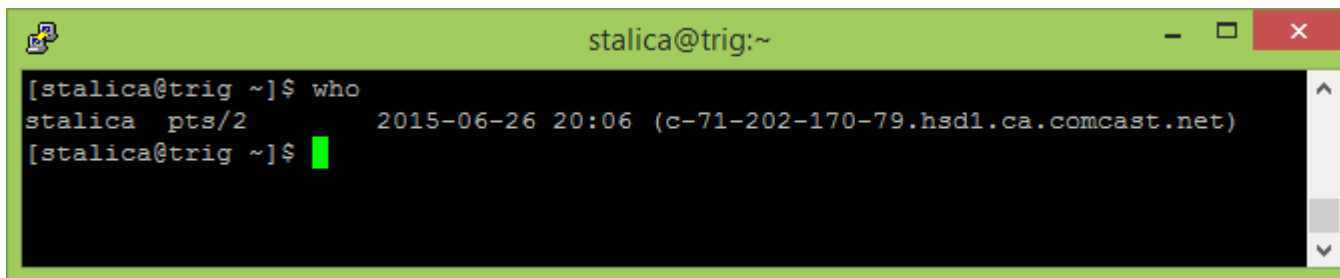
- To learn to use Unix to process files, we'll study the who command.
- We will ask:
 1. What does who do?
 2. How does who work?
 3. Can I write who?

The who command

- Unix commands are things you type at the prompt of a Unix system to do things: ls, cat, who, cp, mkdir, etc
- Commands are just programs written in C
- To add a new command to a Unix system, write a new program and place the executable in a standard system directory such as /bin, /usr/bin, or /usr/local/bin

Question 1: what does who do?

- Tells us who is currently using the system:

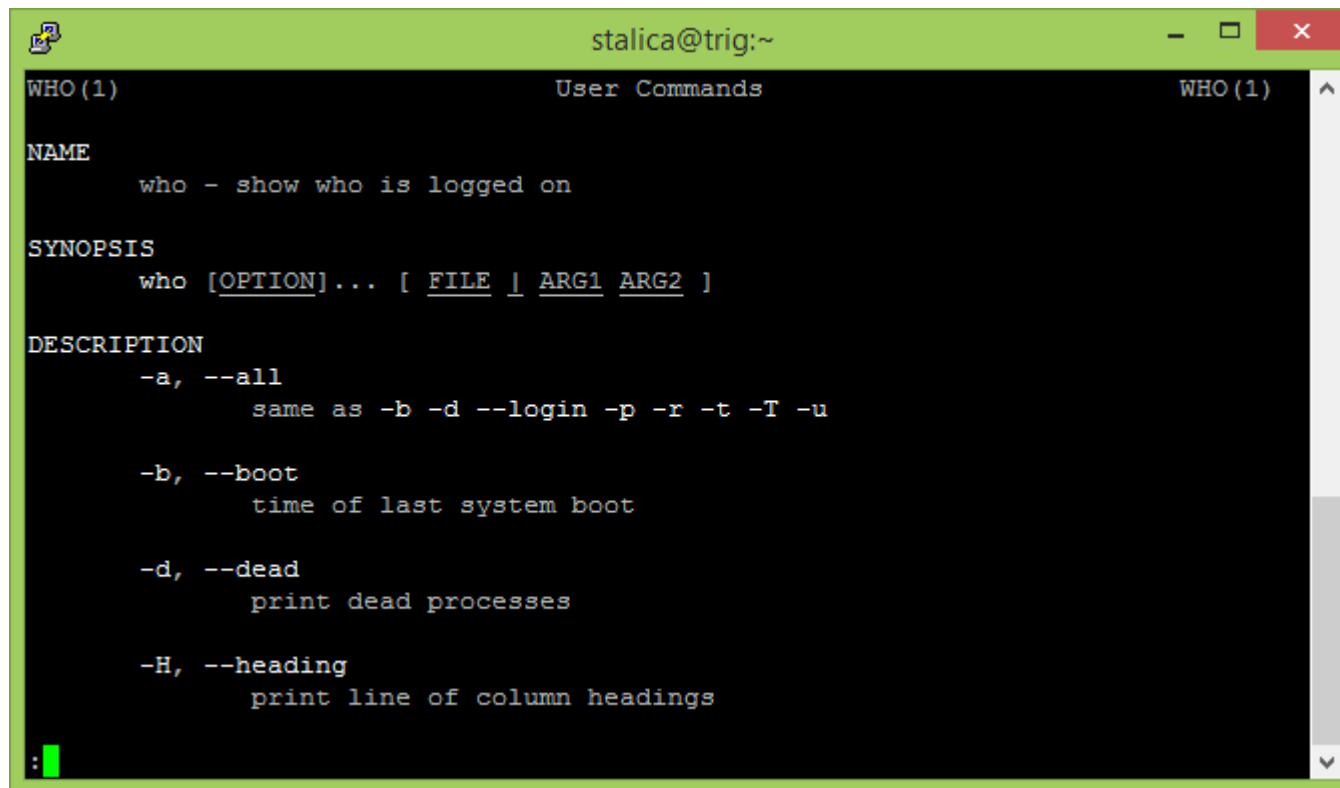


```
stalice@trig:~  
[stalice@trig ~]$ who  
stalice pts/2      2015-06-26 20:06 (c-71-202-170-79.hsd1.ca.comcast.net)  
[stalice@trig ~]$
```

A terminal window with a green title bar containing the text 'stalice@trig:~'. The terminal has a black background with white text. It shows the command 'who' being executed, which returns the output 'stalice pts/2 2015-06-26 20:06 (c-71-202-170-79.hsd1.ca.comcast.net)'. The prompt '[stalice@trig ~]\$' is shown again on the next line with a green cursor.

Question 1: what does who do?

- For more info, check the man page:



A terminal window titled 'stalica@trig:~' displays the man page for the 'who' command. The window has a green title bar and standard window controls. The man page content is as follows:

```
WHO(1)                                User Commands                                WHO(1)
NAME
    who - show who is logged on

SYNOPSIS
    who [OPTION]... [ FILE | ARG1 ARG2 ]

DESCRIPTION
    -a, --all
        same as -b -d --login -p -r -t -T -u

    -b, --boot
        time of last system boot

    -d, --dead
        print dead processes

    -H, --heading
        print line of column headings

:
```

Question 1: what does who do?

- All man pages share the same format
- Top line tells the name of the command and it's section
- Name section contains command name and a brief description.
- Synopsis section shows how to use the command.
- Square brackets [] indicate optional arguments
- Description section details what the command does.

Question 2: How does who do it?

- Let's ask Unix.
- Learn about any Unix command by:
 1. Read the manual
 2. Search the manual
 3. Read the .h files
 4. follow SEE ALSO links
 5. Google it

Question 2: How does who do it?

- From the description section of the who man page, we can see:

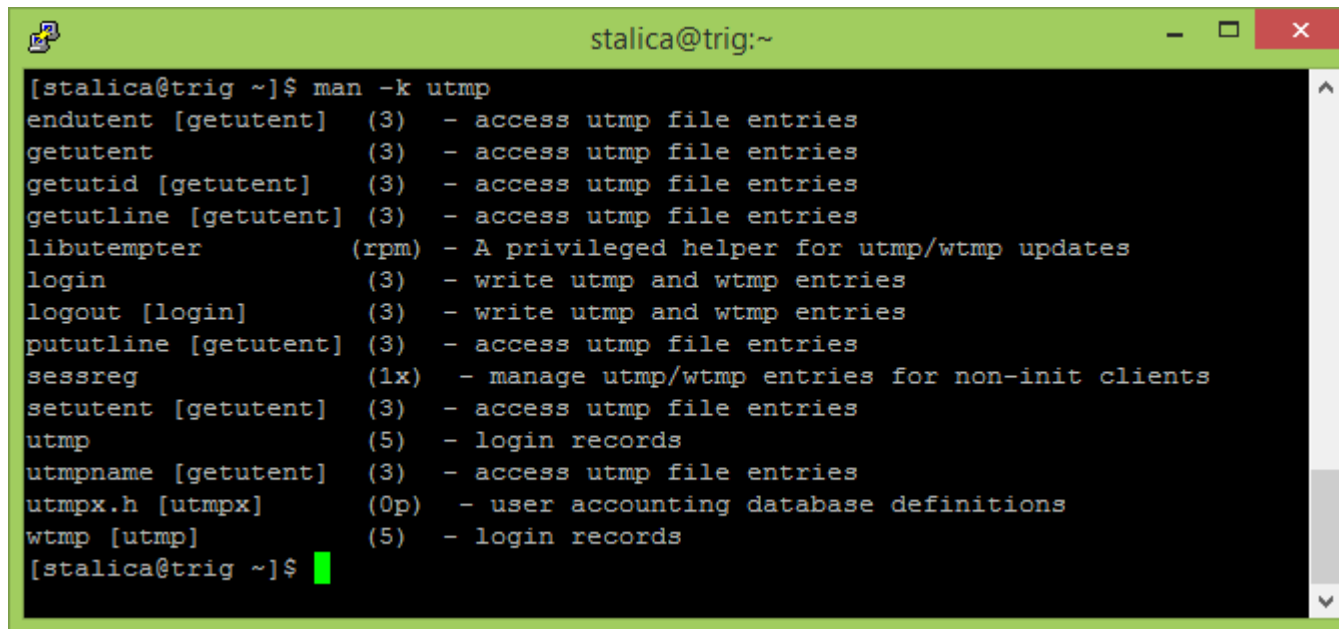


```
stalica@trig:~  
  
--help display this help and exit  
  
--version  
    output version information and exit  
  
If FILE is not specified, use /var/run/utmp. /var/log/wtmp as FILE is  
common. If ARG1 ARG2 given, -m presumed: 'am i' or 'mom likes' are  
usual.  
  
AUTHOR  
:
```

- Interesting, it mentions /var/run/utmp. Let's look for that

Question 2: How does who do it?

- Using `man -k utmp`, we're searching for keyword `utmp`:

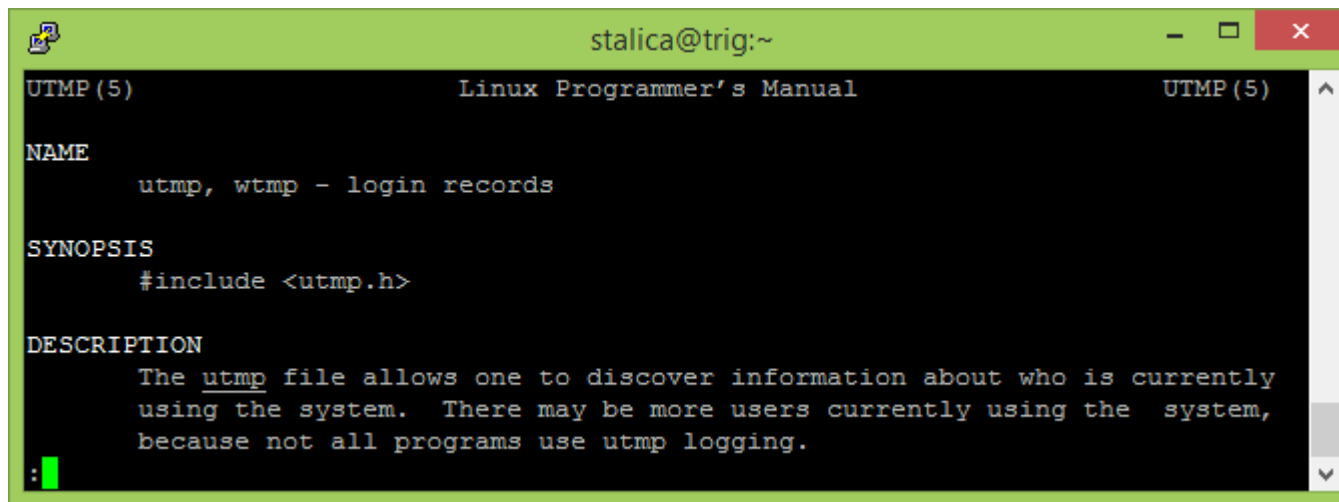


```
stalica@trig:~  
[stalica@trig ~]$ man -k utmp  
endutent [getutent] (3) - access utmp file entries  
getutent (3) - access utmp file entries  
getutid [getutent] (3) - access utmp file entries  
getutline [getutent] (3) - access utmp file entries  
libutempter (rpm) - A privileged helper for utmp/wtmp updates  
login (3) - write utmp and wtmp entries  
logout [login] (3) - write utmp and wtmp entries  
pututline [getutent] (3) - access utmp file entries  
sessreg (1x) - manage utmp/wtmp entries for non-init clients  
setutent [getutent] (3) - access utmp file entries  
utmp (5) - login records  
utmpname [getutent] (3) - access utmp file entries  
utmpx.h [utmpx] (0p) - user accounting database definitions  
wtmp [utmp] (5) - login records  
[stalica@trig ~]$
```

- Hmm, `utmp (5)` – login records, let's look at that

Question 2: How does who do it?

- Let's type, `man 5 utmp`



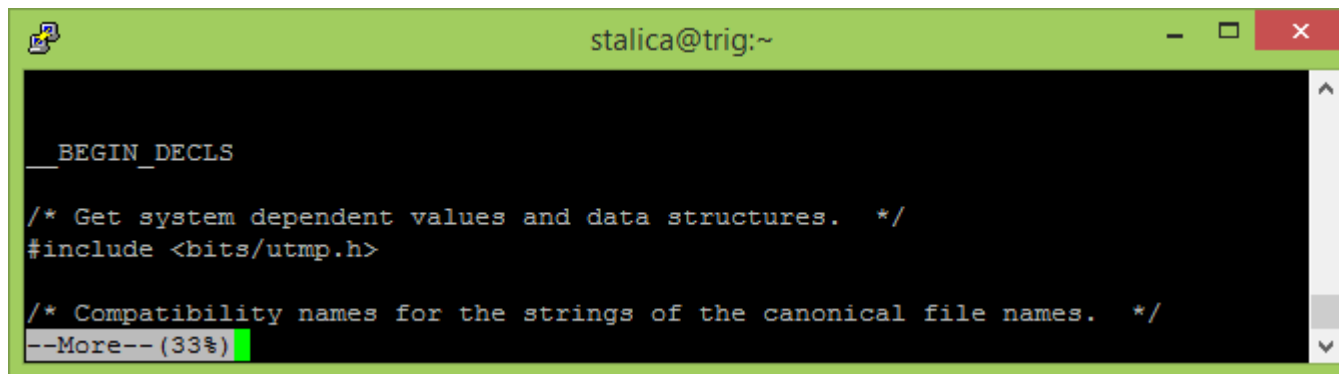
A terminal window titled 'stalice@trig:~' with a green title bar. The window displays the manual page for 'UTMP (5)' from the 'Linux Programmer's Manual'. The content is as follows:

```
UTMP (5)                                Linux Programmer's Manual          UTMP (5)
NAME
    utmp, wtmp - login records
SYNOPSIS
    #include <utmp.h>
DESCRIPTION
    The utmp file allows one to discover information about who is currently
    using the system. There may be more users currently using the system,
    because not all programs use utmp logging.
:
```

- Bingo! Who must use `utmp`, and it's found in `utmp.h`

Question 2: How does who do it?

- Let's read the .h file: `more /usr/include/utmp.h:`

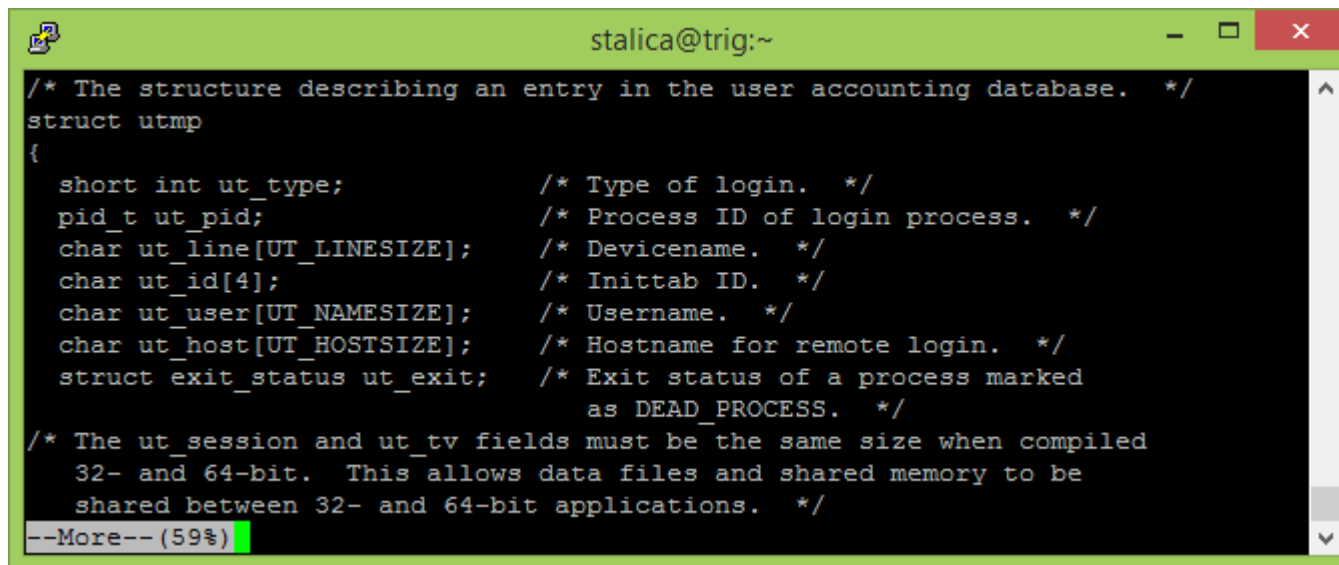
A terminal window with a green title bar containing the text 'stalica@trig:~'. The terminal has a black background with white text. It shows the beginning of the file /usr/include/utmp.h, including the line '__BEGIN_DECLS', a comment '/* Get system dependent values and data structures. */', the include directive '#include <bits/utmp.h>', another comment '/* Compatibility names for the strings of the canonical file names. */', and a prompt '--More-- (33%)' on the last line, which is highlighted in green. A vertical scrollbar is visible on the right side of the terminal window.

```
stalica@trig:~  
__BEGIN_DECLS  
  
/* Get system dependent values and data structures.  */  
#include <bits/utmp.h>  
  
/* Compatibility names for the strings of the canonical file names.  */  
--More-- (33%)
```

- I scroll down and see this, so let's look in there....

Question 2: How does who do it?

- more /usr/include/bits/utmp.h

A terminal window with a green title bar showing the command 'more /usr/include/bits/utmp.h'. The window displays the C code for the 'struct utmp' structure, which is used for user accounting. The code includes comments for each field: 'ut_type' (Type of login), 'ut_pid' (Process ID of login process), 'ut_line' (Devicename), 'ut_id' (Inittab ID), 'ut_user' (Username), 'ut_host' (Hostname for remote login), and 'ut_exit' (Exit status of a process marked as DEAD_PROCESS). It also includes a comment about the 'ut_session' and 'ut_tv' fields being the same size for 32-bit and 64-bit applications. The terminal shows the first part of the file, with a green cursor at the bottom of the first line of code.

```
stalica@trig:~  
/* The structure describing an entry in the user accounting database. */  
struct utmp  
{  
    short int ut_type;          /* Type of login. */  
    pid_t ut_pid;              /* Process ID of login process. */  
    char ut_line[UT_LINESIZE]; /* Devicename. */  
    char ut_id[4];             /* Inittab ID. */  
    char ut_user[UT_NAMESIZE]; /* Username. */  
    char ut_host[UT_HOSTSIZE]; /* Hostname for remote login. */  
    struct exit_status ut_exit; /* Exit status of a process marked  
                                as DEAD_PROCESS. */  
    /* The ut_session and ut_tv fields must be the same size when compiled  
       32- and 64-bit. This allows data files and shared memory to be  
       shared between 32- and 64-bit applications. */  
    --More-- (59%)
```

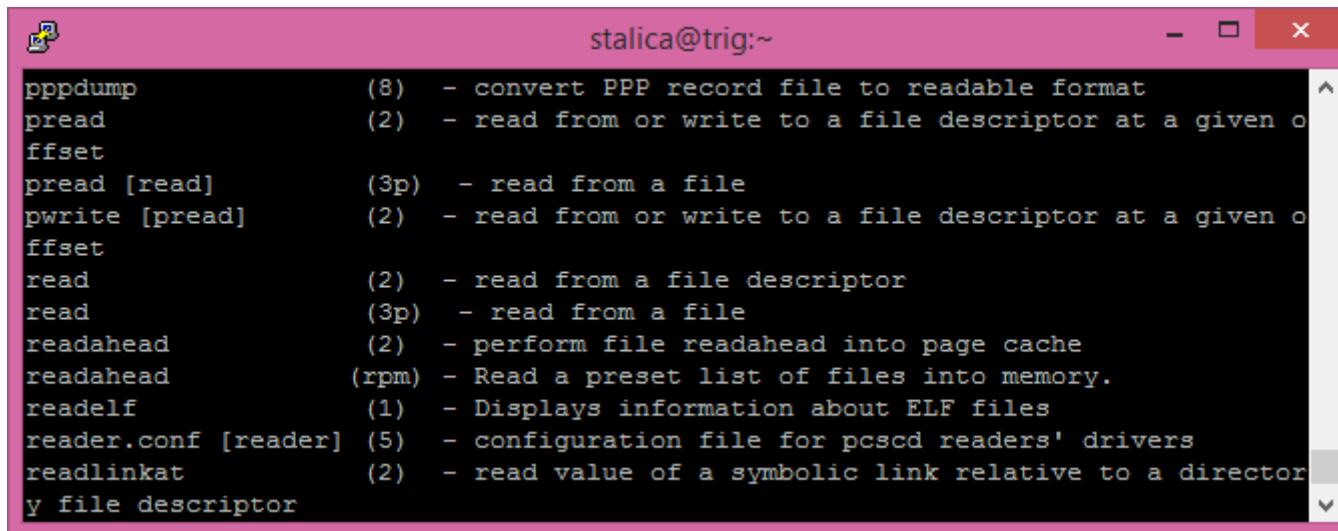
- Awesome. Found it. Now we know how who works.

Question 3: Can we write it?

- We need to do a couple things: read structs from a file, display the information in the struct.
- How do we read structs from a file?
- Read the manual: `man -k file` shows us a LOT of stuff.
- Let's narrow it down, `man -k file | grep read`

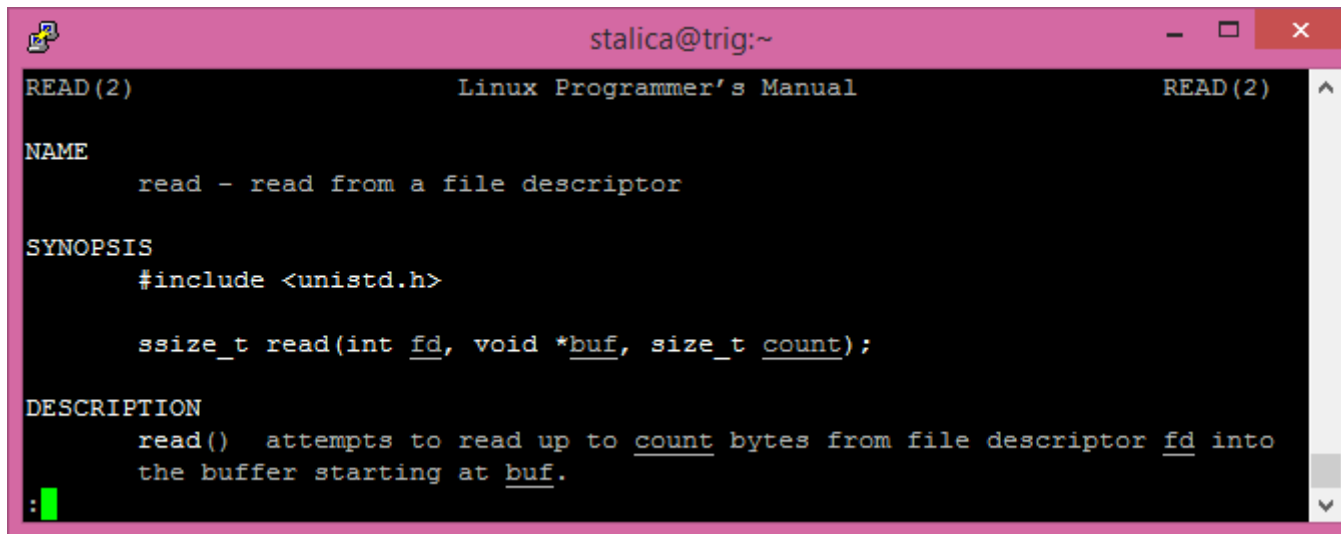
Question 3: Can we write it?

- Read looks promising...

A terminal window with a pink title bar containing the text 'stalica@trig:~'. The terminal has a black background with white text. It displays a list of commands and their descriptions, with a vertical scrollbar on the right side.

```
pppdump      (8) - convert PPP record file to readable format
pread        (2) - read from or write to a file descriptor at a given o
ffset
pread [read] (3p) - read from a file
pwrite [pread] (2) - read from or write to a file descriptor at a given o
ffset
read         (2) - read from a file descriptor
read         (3p) - read from a file
readahead    (2) - perform file readahead into page cache
readahead    (rpm) - Read a preset list of files into memory.
readelf      (1) - Displays information about ELF files
reader.conf [reader] (5) - configuration file for pcsd readers' drivers
readlinkat   (2) - read value of a symbolic link relative to a director
y file descriptor
```

Question 3: Can we write it?



```
stolica@trig:~  
READ(2)                                Linux Programmer's Manual                                READ(2)  
  
NAME  
    read - read from a file descriptor  
  
SYNOPSIS  
    #include <unistd.h>  
  
    ssize_t read(int fd, void *buf, size_t count);  
  
DESCRIPTION  
    read() attempts to read up to count bytes from file descriptor fd into  
    the buffer starting at buf.  
:
```

- So, read is a system call that lets us read n bytes from a file. We can use sizeof to help us out.

Question 3: Can we write it?

- Later on in the man page, system call open is referenced, and open's man page references the system call close.
- We have all we need to do the job.

Question 3: Can we write it?

- Basic usage of the open system call:

open

PURPOSE:	Creates a connection to a file	
INCLUDE:	#include <fcntl.h>	
USAGE:	int fd = open(char* name, int how)	
ARGS:	name	name of file
	how	O_RDONLY, O_WRONLY, or O_RDWR
RETURNS:	-1	on error
	int	on success

Question 3: Can we write it?

- Opening a file is a kernel service.
- Unix allows several processes to open the same file.
- If the file opens successfully, the kernel returns a small positive integer known as a ***file descriptor*** which is used to identify the connection.
- The file descriptor is used for all operations involving the file.

Question 3: Can we write it?

- We read from the file using the file descriptor.

read		
PURPOSE	Transfer up to qty bytes from fd to buf	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t numread = read(int fd, void* buf, ssize_t qty)	
ARGS	fd	source of data
	buf	destination of data
	qty	number of bytes to transfer
RETURNS	-1	on error
	numread	on success

Question 3: Can we write it?

- Need to close a file when we're done with it, here's the call:

close

PURPOSE	Close a file
----------------	--------------

INCLUDE	#include <unistd.h>
----------------	---------------------

USAGE	int result = close(int fd)
--------------	------------------------------

ARGS	fd file descriptor
-------------	-------------------------

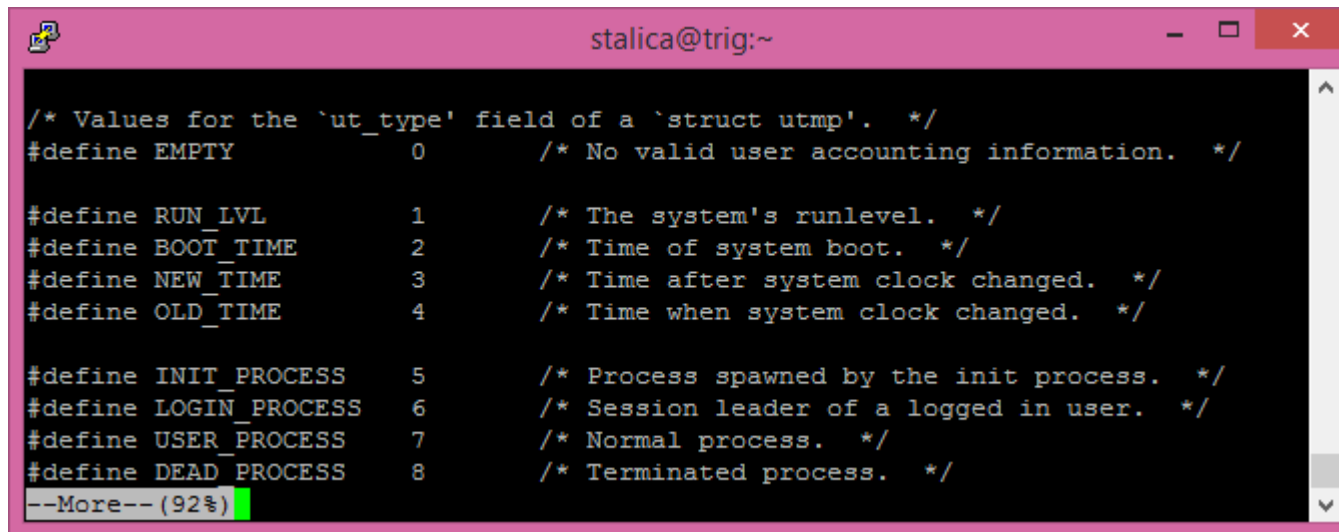
RETURNS	-1 on error
	0 on success

Question 3: Can we write it?

- Let's walk through who1.c, the first version, paying specific attention to the show_info function.
- How closely does who1.c match the real who?
- Not close enough, we need to suppress blank records and get log-in times correct.

Question 3: Can we write it?

- Doing some further digging around in utmp.h, we find:

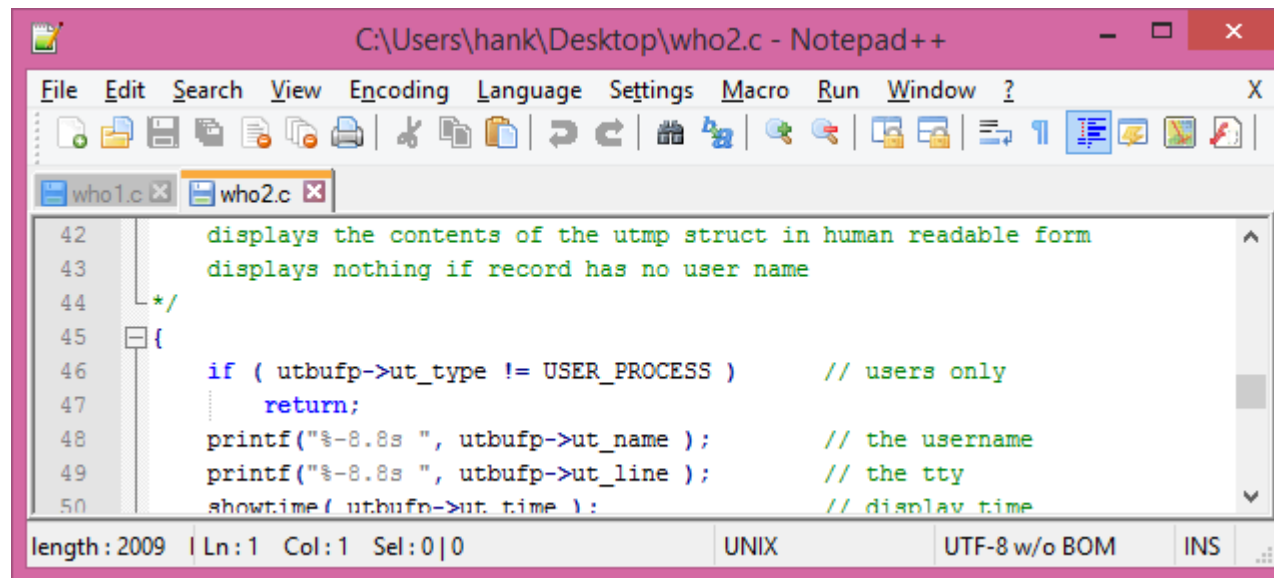
A terminal window with a pink title bar containing the text 'stalica@trig:~'. The terminal has a black background with white text. It displays the definitions for the 'ut_type' field from the utmp.h header file. The definitions are listed as follows: 0 for EMPTY (No valid user accounting information), 1 for RUN_LVL (The system's runlevel), 2 for BOOT_TIME (Time of system boot), 3 for NEW_TIME (Time after system clock changed), 4 for OLD_TIME (Time when system clock changed), 5 for INIT_PROCESS (Process spawned by the init process), 6 for LOGIN_PROCESS (Session leader of a logged in user), 7 for USER_PROCESS (Normal process), and 8 for DEAD_PROCESS (Terminated process). At the bottom of the terminal, the text '--More-- (92%)' is visible, with a green cursor positioned at the end of the line.

```
stalica@trig:~  
/* Values for the `ut_type' field of a `struct utmp'. */  
#define EMPTY          0          /* No valid user accounting information. */  
  
#define RUN_LVL        1          /* The system's runlevel. */  
#define BOOT_TIME      2          /* Time of system boot. */  
#define NEW_TIME       3          /* Time after system clock changed. */  
#define OLD_TIME       4          /* Time when system clock changed. */  
  
#define INIT_PROCESS   5          /* Process spawned by the init process. */  
#define LOGIN_PROCESS  6          /* Session leader of a logged in user. */  
#define USER_PROCESS   7          /* Normal process. */  
#define DEAD_PROCESS   8          /* Terminated process. */  
--More-- (92%)
```

- Type 7 looks like something we can use

Question 3: Can we write it?

- So let's make fix one:



The screenshot shows a Notepad++ window titled "C:\Users\hank\Desktop\who2.c - Notepad++". The window has a menu bar with File, Edit, Search, View, Encoding, Language, Settings, Macro, Run, and Window. Below the menu is a toolbar with various icons. The main text area shows C code for a function that displays utmp struct contents. The code is as follows:

```
42      displays the contents of the utmp struct in human readable form
43      displays nothing if record has no user name
44  */
45  {
46      if ( utbufp->ut_type != USER_PROCESS )      // users only
47          return;
48      printf("%-8.8s ", utbufp->ut_name );          // the username
49      printf("%-8.8s ", utbufp->ut_line );          // the tty
50      showtime( utbufp->ut_time );                  // display time
```

The status bar at the bottom indicates "length: 2009 | Ln: 1 Col: 1 Sel: 0 | 0", "UNIX", "UTF-8 w/o BOM", and "INS".

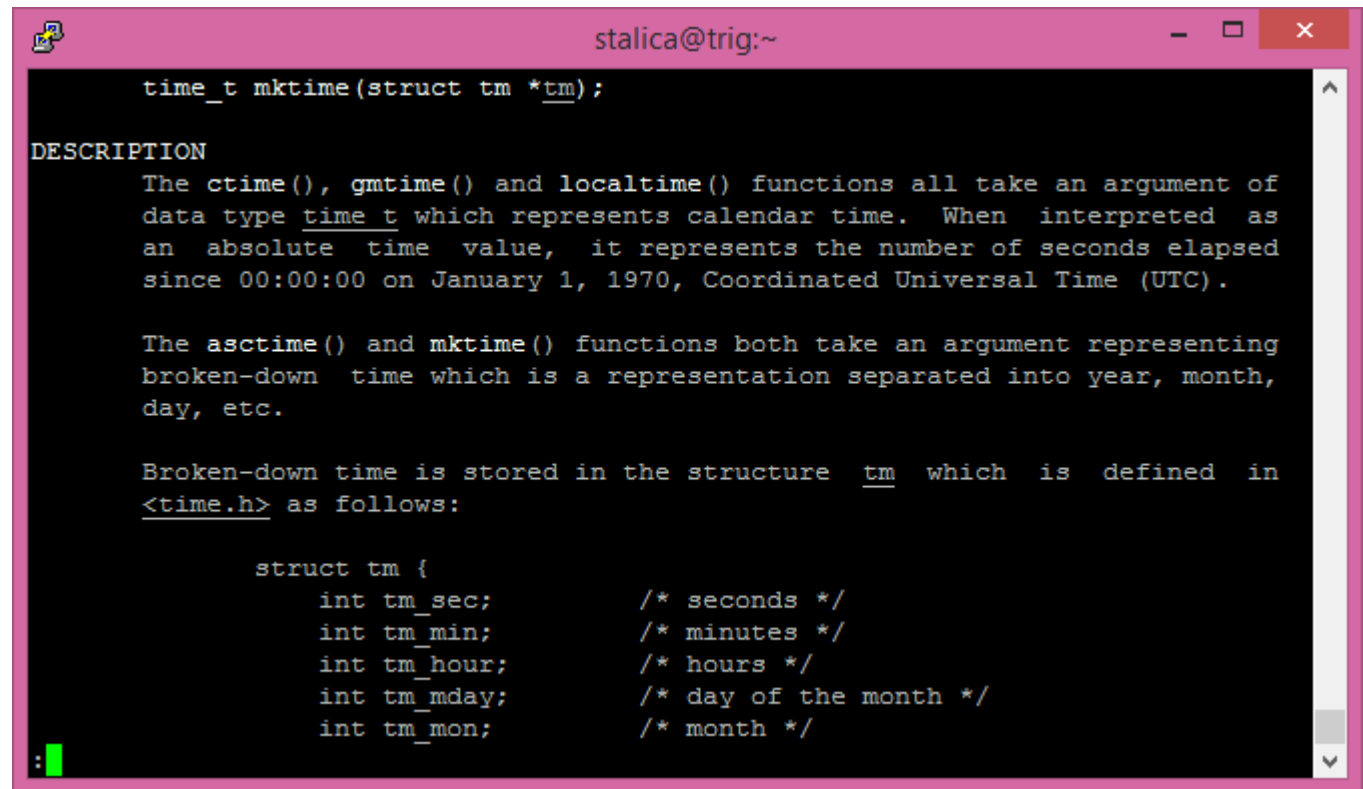
- Lines 46, 47 fix this problem.

Question 3: Can we write it?

- The next problem is fixing that time display `man -k time` which leads us to
- `/usr/include/time.h` has information we can use
- Unix has a data type it uses to store time, `time_t`.
- It's an integer that stores all the seconds since midnight, Jan 1, 1970.
- `ut_time` in `utmp` represent the login-time as seconds since the beginning of *The Epoch*.

Question 3: Can we write it?

- Need to convert that to something readable. Let's use a function to do that: `ctime`
- Man 3 `ctime`:

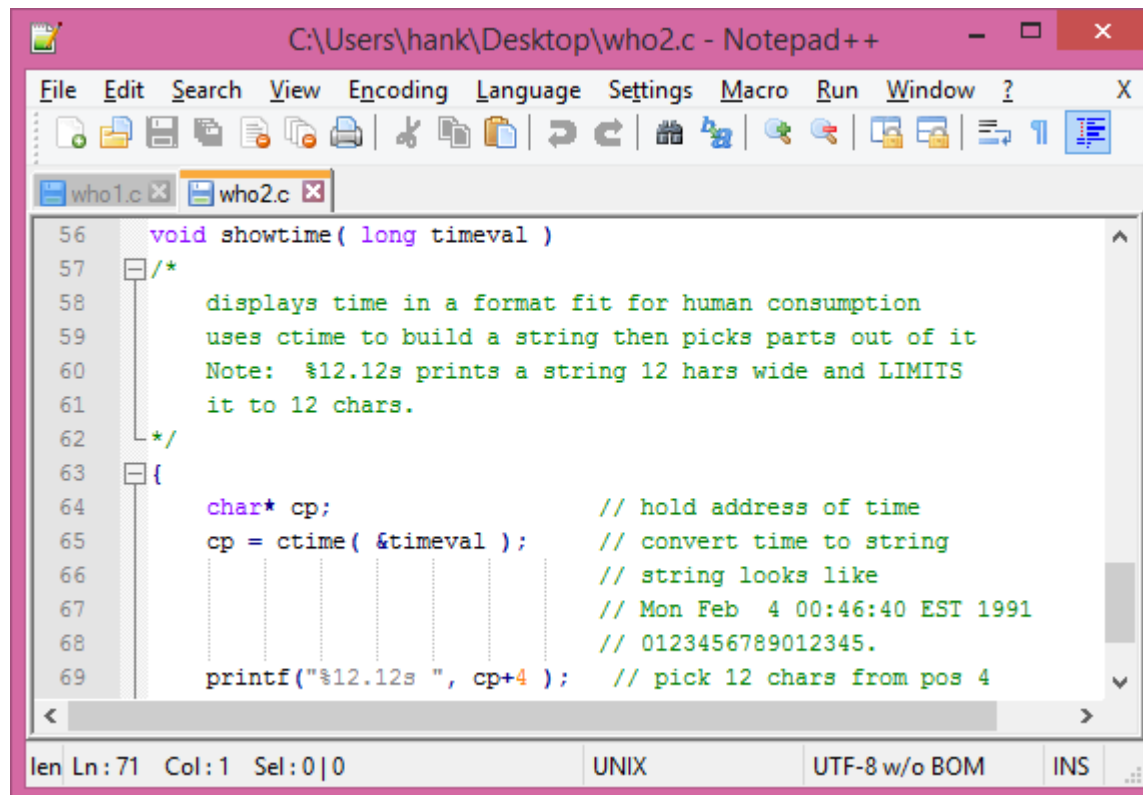


```
stalica@trig:~  
time_t mktime(struct tm *tm);  
  
DESCRIPTION  
The ctime(), gmtime() and localtime() functions all take an argument of  
data type time_t which represents calendar time. When interpreted as  
an absolute time value, it represents the number of seconds elapsed  
since 00:00:00 on January 1, 1970, Coordinated Universal Time (UTC).  
  
The asctime() and mktime() functions both take an argument representing  
broken-down time which is a representation separated into year, month,  
day, etc.  
  
Broken-down time is stored in the structure tm which is defined in  
<time.h> as follows:  
  
    struct tm {  
        int tm_sec;           /* seconds */  
        int tm_min;           /* minutes */  
        int tm_hour;          /* hours */  
        int tm_mday;          /* day of the month */  
        int tm_mon;           /* month */
```

Question 3: Can we write it?

- Looks like just what we need. The function takes a `time_t` pointer and returns a pointer to a string that looks something like: `Wed Jun 30 21:49:08 1993\n`
- Let's put it all together and look at `who2.c`, our finished version.

Question 3: Can we write it?



The screenshot shows a Notepad++ window titled "C:\Users\hank\Desktop\who2.c - Notepad++". The window contains two tabs: "who1.c" and "who2.c". The "who2.c" tab is active and displays the following C code:

```
56 void showtime( long timeval )
57 /*
58     displays time in a format fit for human consumption
59     uses ctime to build a string then picks parts out of it
60     Note: %12.12s prints a string 12 chars wide and LIMITS
61     it to 12 chars.
62 */
63 {
64     char* cp;                // hold address of time
65     cp = ctime( &timeval ); // convert time to string
66                               // string looks like
67                               // Mon Feb  4 00:46:40 EST 1991
68                               // 0123456789012345.
69     printf("%12.12s ", cp+4 ); // pick 12 chars from pos 4
```

The status bar at the bottom of the window shows "len Ln: 71 Col: 1 Sel: 0|0", "UNIX", "UTF-8 w/o BOM", and "INS".

Project Two: Writing cp (read and write)

Writing cp

- Let's follow a similar procedure, and write our own cp command.
- Remember the approach,
 1. What does cp do?
 2. How does cp Create and Write?
 3. Can we write our own cp?

What does cp do?

- The cp system command makes a copy of a file.
- Typical usage: `cp source-file target-file`
- If there is no target file, cp creates it. If there is, cp replaces it.

How does cp do it?

- We can create or rewrite a file using the creat system call.

creat

PURPOSE	Create or zero a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = creat(char* fn, mode_t mode)	
ARGS	fn:	the name of the file
	mode:	access permission
RETURNS	-1	on error
	fd	on success

How does cp do it?

- An example,
 `fd = creat("addressbook", 0644);`
- This creates or truncates a file named addressbook.
- If the file doesn't exist, the permissions are set to
 `rw--r--r--` which we'll explain in Chapter 3.
- `fd` represents a file open for writing only

How does cp do it?

- Let's send data to an open file using the write system call:

write

PURPOSE	send data from memory to a file
----------------	---------------------------------

INCLUDE	#include <unistd.h>
----------------	---------------------

USAGE	ssize_t result = write(int fd, void* buf, size_t amt)
--------------	---

ARGS	fd	a file descriptor
	buf	an array
	amt	how many bytes to write

RETURNS	-1	on error
	num bytes written	on success

How does cp do it?

- So write copies data from process memory to a file.
- If there's a problem, the kernel returns -1.
- Otherwise, returns number of bytes written.

Can we write cp?

- Here's a general outline for cp:

```
open sourcefile for reading
open copyfile for writing
while not end of file do
    read from sourcefile to buffer
    write from buffer to copyfile
close sourcefile
close copyfile
```


Can we write cp?

- Consider what's going to happen:
- The files are on the disk.
- Process is living in user space.
- A buffer is a chunk of memory in the process.
- Process has two file descriptors.
- Bytes are read from the original file into the buffer then to the copy
- Let's look at an implementation, cp1.c. (see sample code)

More efficient file i/o: buffering

- Let's ask a new question: *How can I make this run better?*
- `cp1` contains a symbol `BUFFERSIZE`.
- Defines size of the array that holds the bytes being read/written.
- Does the size of the buffer matter?

Does buffer size matter?

- Absolutely. Consider copying a file 2500 bytes long. Then,

Filesize = 2500 bytes

If buffer = 100 bytes, then
copy requires 25 read() and 25 write() calls

If buffer = 1000 bytes, then
copy requires 3 read() and 3 write() calls

Does buffer size matter?

- Going from 100 \rightarrow 1000 means we go from 50 calls to 6.
- HUGE difference.
- System calls take time. Programs that make a lot of them run slower and take cpu time away from other users.

Why System Calls Consume Time

- Consider the cp1 program.
- In our system's memory, the cp1 process is in *user space*, and the kernel is in *system space* : two distinct chunks of memory.
- cp1 wants to do a read, but the code to do the reading is actually in the kernel. So, control jumps from the code in cp1 in user space to the kernel code in system space.
- CPU runs the code, this takes time to transfer the data.

Why System Calls Consume Time

- Additionally, jumping into and out of the kernel takes time.
- When the kernel code is running, the CPU is running in *supervisor* mode with its own stack and memory.
- When user code is running, the CPU is running in *user* mode, with its own stack and memory.
- The kernel needs special access to resources that user code must not have.
- The more system calls that are made, the more time is spent switching between these modes. This is *expensive*.

Does this mean who2.c is inefficient?

- Yes! One system call for reading each utmp record makes as much sense as buying a carton of eggs, one egg at a time.
- Better idea: read a bunch of records at once, then process the records in local storage one by one.
- Consider that egg carton, and some pseudocode:

Does this mean who2.c is inefficient?

```
getegg() {  
    if ( eggs_left_in_carton == 0 ) {  
        refill carton at store  
        if ( eggs_at_store == 0 )  
            return EndOfEggs  
        eggs_left_in_carton = 12  
    }  
    eggs_left_in_carton--:  
    return one egg;  
}
```


Does this mean who2.c is inefficient?

- Each call to `getegg` fetches one egg, but not from the store
- Only when the carton is empty does the function go to the store.
- If you look at a lot of Unix functions, they are implemented using similar logic.
- Adding buffering to `who2.c` is left as an exercise.

Buffering and the Kernel

- If buffering is so great, why doesn't the kernel do it? *It does.*
- Switching between modes takes time, but reading disks takes forever in comparison.
- To save time, the kernel keeps copies of disk blocks in memory.
- Disks are collection of blocks of data, similar to how the utmp file is a collection of log-in records.
- Kernel copies blocks from the disk into kernel buffers.

Buffering and the Kernel

- When a process requests file data, the kernel copies it from its own buffer to the process buffer, not from the disk to user space.
- What happens if the data isn't in the kernel buffer?
- Kernel adds it to the shopping list, and suspends the requesting process.
- Kernel lets other processes run in the mean time.

Buffering and the Kernel

- Later, the kernel moves the requested data from disk into the kernel buffer, and can then copy the data into the buffer in user space and wake the sleeping requesting process.
- So, `read()` causes the kernel to copy data from kernel buffer into the process buffer. `write()` causes data to be copied from the process to a kernel buffer.
- The kernel is then free to copy data to disk when it gets around to it.

Buffering and the Kernel

- It's possible those kernel buffers don't get written to disk.
- Make sure you shut down your computer correctly!
- Consequences of Kernel Buffering:

Faster "disk" I/O

Optimized disk writes

Need to write buffers to disk before shutdown

Reading and Writing a File

- Recall the read/write pointers from C++? Unix has a system call that mimics this behavior, lseek.
- lseek allows us to set the file pointer to a specific offset within the file.
- We can then read from that offset, or overwrite data beginning at that offset.
- Let's look at the call:

Reading and Writing a File

- **lseek**

PURPOSE	set file pointer to specified offset in file						
INCLUDE	<code>#include <sys/types.h></code> <code>#include <unistd.h></code>						
USAGE	<code>off_t oldpos = lseek(int fd, off_t dist, int base)</code>						
ARGS	<table><tr><td>fd:</td><td>file descriptor</td></tr><tr><td>dist:</td><td>a distance in bytes</td></tr><tr><td>base:</td><td><code>SEEK_SET</code> ==> start of file <code>SEEK_CUR</code> ==> current position <code>SEEK_END</code> ==> end of file</td></tr></table>	fd:	file descriptor	dist:	a distance in bytes	base:	<code>SEEK_SET</code> ==> start of file <code>SEEK_CUR</code> ==> current position <code>SEEK_END</code> ==> end of file
fd:	file descriptor						
dist:	a distance in bytes						
base:	<code>SEEK_SET</code> ==> start of file <code>SEEK_CUR</code> ==> current position <code>SEEK_END</code> ==> end of file						
RETURNS	-1 on error or the previous position in the file						

Reading and Writing a File

- lseek sets the current position for open file *fd* to the position defined by the pair *dist* and *base*.

- Example:

```
lseek( fd, -(sizeof( struct utmp ) ), SEEK_CUR);
```

That's gonna move the current position a distance `sizeof(struct utmp)` bytes *before* the current position.

- Notice: `lseek(fd, 0, SEEK_CUR)` returns the current position.

Reading and Writing a File

- Let's see an example, lseek.c.

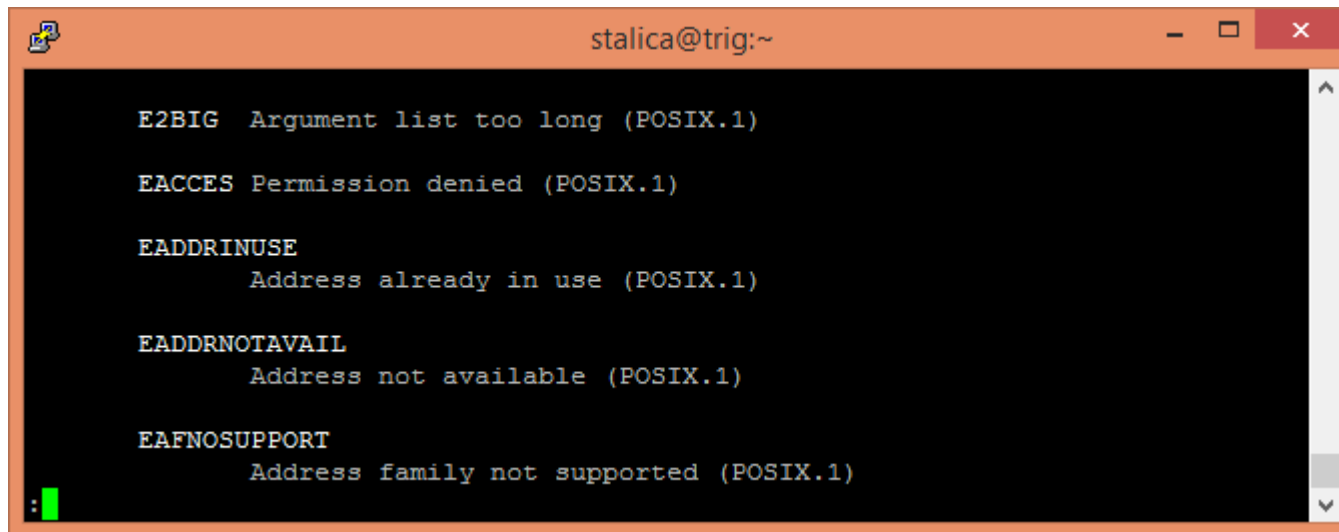
What To Do with System-Call Errors

errno

- Every system call has its own set of errors.
- Consider `open()`. The file might not exist, you might not have permissions, you might have too many files open.
- How to identify what went wrong: `errno`
- The kernel will tell your program the error cause by storing an error code in a global variable called `errno`.
- Every program can access this variable.

errno

- The manpage for `errno(3)` include error-code symbols:



A terminal window titled 'stalica@trig:~' with a black background and orange border. It displays a list of error codes and their descriptions from the `errno(3)` manpage. The text is as follows:

```
E2BIG  Argument list too long (POSIX.1)

EACCES  Permission denied (POSIX.1)

EADDRINUSE
        Address already in use (POSIX.1)

EADDRNOTAVAIL
        Address not available (POSIX.1)

EAFNOSUPPORT
        Address family not supported (POSIX.1)
```

A green cursor is visible at the bottom left of the terminal window.

errno

- You can use these symbols to figure out what went wrong.
- For example:

```
1 #include <errno.h>
2 #include <fcntl.h>
3
4 int main()
5 {
6     int fd = open( "file", O_RDONLY );
7
8     if ( fd == -1 )
9     {
10         int errsv = errno;
11
12         if ( errsv == ENOENT )
13             printf("No such file\n");
14         else if ( errsv == EACCESS )
15             printf("You do not have permission to open file\n");
16     }
```

errno

- To access errno on older systems, an extern call is required:
`extern int errno;`
- On modern systems and versions of the C library, this doesn't work. Instead, include a header file:
`#include <errno.h>`
- Be sure to save the error message before making another system call! Otherwise, you risk overwriting the previous error number.

Reporting Errors: perror(3)

- You can print an error message describing the error doing as above.
- Often, it's easier to use a built-in function perror(string).
- This function looks up the error code and prints, to standard error, the string you pass it and a descriptive message.

Reporting Errors: perror(3)

- Here's a revised version of the previous sample:

```
1 #include <errno.h>
2 #include <fcntl.h>
3 #include <stdio.h>
4
5 int main()
6 {
7     int fd = open( "file", O_RDONLY );
8
9     if ( fd == -1 )
10    {
11        perror("Can't open file\n");
12        exit(1);
13    }
```


Summary

- who command reports the list of current users by reading a system log file.
- Unix systems store data in files. Unix manipulates this data using six system calls:

open(filename, how)
creat(filename, mode)
read(fd, buffer, amt)
write(fd, buffer, amt)
lseek(fd, distance, base)
close(fd)

Summary

- A process reads/writes data through *file descriptors* which identify a connection between a process and a file.
- Each time a process makes a system call, the computer changes from user mode to kernel mode and executes kernel code. Programs are more efficient when this is minimized.
- Programs can reduce the number of system calls when reading/writing data by using buffers and calling the kernel when they are full/empty.

Summary

- Unix kernel uses buffers inside kernel memory to minimize data transfers.
- Unix stores times as the number of seconds since beginning of Unix time.
- When Unix system calls have an error, the system sets `errno` to an error code and returns the value `-1`. That error code can be used to diagnose the error and take action.
- Use man pages to find out about commands.
- Header files contain definitions of structures, symbolic constant values, and function prototypes used to build system tools.