

CH08: Processes and Programs

Studying sh

Objectives

Ideas and Skills

- What a Unix shell does
- The Unix model of a process
- How to run a program
- How to create a process
- How parent and child processes communicate

System Calls, Functions, Commands

- fork
- exec
- wait
- exit
- sh
- ps

8.1: Processes = Programs In Action

Processes = Programs in Action

- A program is a sequence of machine-language instruction stored in a file.
- Running a program means loading that list of instructions into memory, then the CPU executes the instructions, one by one.
- In Unix, an **executable program** is a list of machine language and data.

Processes = Programs in Action

- And a **process** is the memory space and settings the program uses to run.




FIGURE 8.1

Processes are programs in action.

Processes = Programs in Action

- Data and programs are stored in files on disk.
- Programs run in processes.
- Let's learn more by playing with ps

8.2: Learning About Processes with *ps*

Learning about processes with ps

- Processes live in user space – the part of memory where running programs and their data live.




FIGURE 8.2
The `ps` command lists current processes

Learning about processes with ps

- Let's find out more using *ps*, which lists current processes:

```
hstalica@cat:~$ ps
  PID TTY          TIME CMD
 3037 pts/2        00:00:00 bash
 3168 pts/2        00:00:00 ps
hstalica@cat:~$
```

Learning about processes with ps

- Here, we see two processes running: bash (a shell) and ps.
- Each process has it's own id number, called a PID.
- Each process is connected to a terminal, here it's /dev/pts/2.
- Each has an elapsed time.

Learning about processes with ps

- ps has an -a option which shows more processes, including ones run by other users.
- -a doesn't show shells.

```
$ ps -a
  PID TTY          TIME CMD
 1779 pts/0        00:00:13 gv
 1780 pts/0        00:00:07 gs
 1781 pts/0        00:00:01 vi
 2013 pts/2        00:00:23 xpaint
 2017 pts/2        00:00:02 mail
 2018 pts/1        00:00:00 ps
```

Learning about processes with ps

- -l option prints more information:

```
$ ps -la
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
000	S	504	1779	1731	0	69	0	-	1086	do_sel	pts/0	00:00:13	gv
000	S	504	1780	1779	0	69	0	-	2309	do_sel	pts/0	00:00:07	gs
000	S	504	1781	1731	0	72	0	-	1320	do_sel	pts/0	00:00:01	vi
000	S	519	2013	1993	0	69	19	-	1300	do_sel	pts/2	00:00:23	xpain
000	S	519	2017	1993	0	69	0	-	363	read_c	pts/2	00:00:02	mail
000	R	500	2023	1755	0	79	0	-	750	-	pts/1	00:00:00	ps

Learning about processes with ps

- Column S shows status of each process.
- ps is running, indicated by the R.
- Each process belongs to a user, and their ID is listed.
- Each process has its PID, but we also see the parent's process ID (PPID).

Learning about processes with ps

- PRI and NI indicate priority and niceness levels of a process. The kernel uses them to decide when to run the process.
- A process can increase its niceness, which allows others to go before it.
- The superuser can decrease the niceness level, pushing ahead of it in line.

Learning about processes with ps

- Processes have a size, shown in SZ column which is the amount of memory taken up.
- The size can change as the process runs (dynamic memory allocation)
- WCHAN shows why a process is sleeping.
- do_sel and read_c indicate they are waiting for input – these are addresses in the kernel.

Learning about processes with ps

- ADDR and F are deprecated.
- These options can work differently from system to system. Check man pages.

System Processes

- Additionally, Unix runs it's own processes to perform system tasks:

```
hstalica@cat:~$ ps -ax | head -25
  PID TTY          STAT TIME  COMMAND
    1 ?           Ss   0:02 /sbin/init splash
    2 ?           S    0:00 [kthreadd]
    3 ?           S    0:00 [ksoftirqd/0]
    5 ?          S<   0:00 [kworker/0:0H]
    6 ?           S    0:00 [kworker/u2:0]
    7 ?           S    0:00 [rcu_sched]
    8 ?           S    0:00 [rcu_bh]
    9 ?           S    0:00 [rcuos/0]
   10 ?           S    0:00 [rcuob/0]
   11 ?           S    0:00 [migration/0]
   12 ?           S    0:00 [watchdog/0]
   13 ?          S<   0:00 [khelper]
   14 ?           S    0:00 [kdevtmpfs]
   15 ?          S<   0:00 [netns]
   16 ?          S<   0:00 [perf]
   17 ?           S    0:00 [khungtaskd]
   18 ?          S<   0:00 [writeback]
   19 ?           SN    0:00 [ksmd]
   20 ?           SN    0:00 [khugepaged]
   21 ?          S<   0:00 [crypto]
   22 ?          S<   0:00 [kintegrityd]
   23 ?          S<   0:00 [bioset]
   24 ?          S<   0:00 [kblockd]
   25 ?          S<   0:00 [ata_sff]
hstalica@cat:~$
```

System Processes

- This shows the first 25 of 163 processes running on my system.
- Some are system processes.
- Many don't have a terminal.
- Some manage memory, others manage system logfiles (klogd, syslogd), etc.

System Processes

- A lot can be learned about a system by reading `ps` – `ax` and man pages.
- Using `ps` allows you to see the number and diversity of processes living in a computer.

Process Management and File Management

- Processes have many attributes: belong to a user ID, have a size, a starting time, elapsed running time, priority, and niceness.
- Where are these properties stored?
- The kernel manages processes in memory and files on the disk.
- How similar are they?

Process Management and File Management

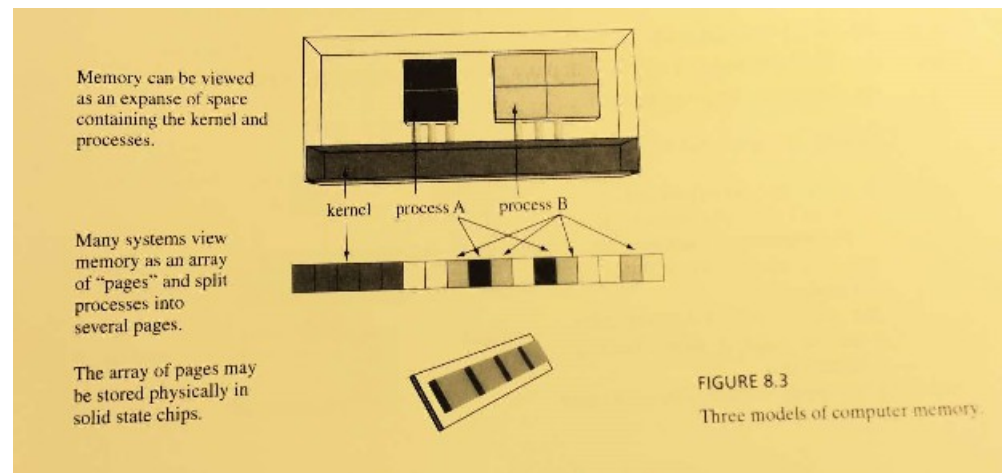
- Files contain data, processes contain executable code.
- Files have attributes, processes have attributes.
- Kernel creates and destroys files, same with processes.
- Kernel stores several processes in memory, and stores several files on disk.

Process Management and File Management

- The kernel has to allocate space and track which processes use which blocks of memory.
- Sound familiar?

Computer Memory and Computer Programs

- A process is an abstract idea, but represents the concrete: bytes in memory.



Computer Memory and Computer Programs

- Unix divides memory into kernel space and user space.
- Processes run in user space.
- Memory is just a big array (like disk blocks).
- A machine with as little as 64 megabytes of memory has an array of ~67 million memory locations.

Computer Memory and Computer Programs

- Some of that is instructions and data that make up the kernel.
- Some of that is instructions and data that compromise processes.
- Processes don't necessarily occupy a single chunk of memory, usually they are divided into smaller chunks. (just like disk files)

Computer Memory and Computer Programs

- Files have an allocation list of disk blocks, processes have a structure to hold the allocation list of memory pages.
- Creating a process is similar to creating a disk file.
- The kernel finds free memory pages to store the instructions and data of the process.
- It then sets up data structures to store memory allocation info and attributes of the process.

8.3 The Shell: A Tool for Process and Program Control

The Shell

- A **shell** is a program that manages processes and runs programs.
- There are many shells available, but the popular ones provide three main functions:
 - (a) Shells run programs
 - (b) Shells manage input and output
 - (c) Shells can be programmed

The Shell

- Consider:

```
hstalica@cat:~$ grep lp /etc/passwd
lp:x:7:7:lp:/var/spool/lpd:/usr/sbin/nologin
hstalica@cat:~$ TZ=PST8PDT ; export TZ ; date ; TZ=EST5EDT
Tue May 17 09:58:10 PDT 2016
hstalica@cat:~$ date
Tue May 17 12:58:20 EDT 2016
hstalica@cat:~$ ls -l /etc > etc.listing
hstalica@cat:~$ NAME=lp
hstalica@cat:~$ if grep $NAME /etc/passwd
> then
>     echo hello | mail $NAME
> fi
lp:x:7:7:lp:/var/spool/lpd:/usr/sbin/nologin
```

Running Programs

- The commands grep, date, ls, echo, and mail are regular programs written in C.
- The shell loads them into memory and runs them.
- A shell can be thought of as a program launcher.

Managing Input and Output

- Using `>`, `<`, and `|`, a user tells the shell to attach input and output of processes to disks files or other processes.


Programming

- The shell is a programming language with variables and flow control.
- TZ is set to a string representing a time-zone.
- That value is used by the date command.
- An *if...then* statement is used.
- If the grep command succeeds in searching for "lp" in /etc/passwd, shell executes a command.

8.4: How The Shell Runs Programs

How the Shell Runs Programs

- You see a prompt, you type a command, the shell runs it, you see the prompt again. What's going on?



How the Shell Runs Programs

- A shell follows the following steps: (fig 8.4)
 - A. The user types a.out
 - B. The shell creates a new process to run the program.
 - C. The shell loads the program from disk into the process.
 - D. The program runs in its process until done
- This is the main loop of the shell.

The Main Loop of a Shell

- The shell consists of the following loop:

```
while( ! end_of_input )  
    get command  
    execute command  
    wait for command to finish
```

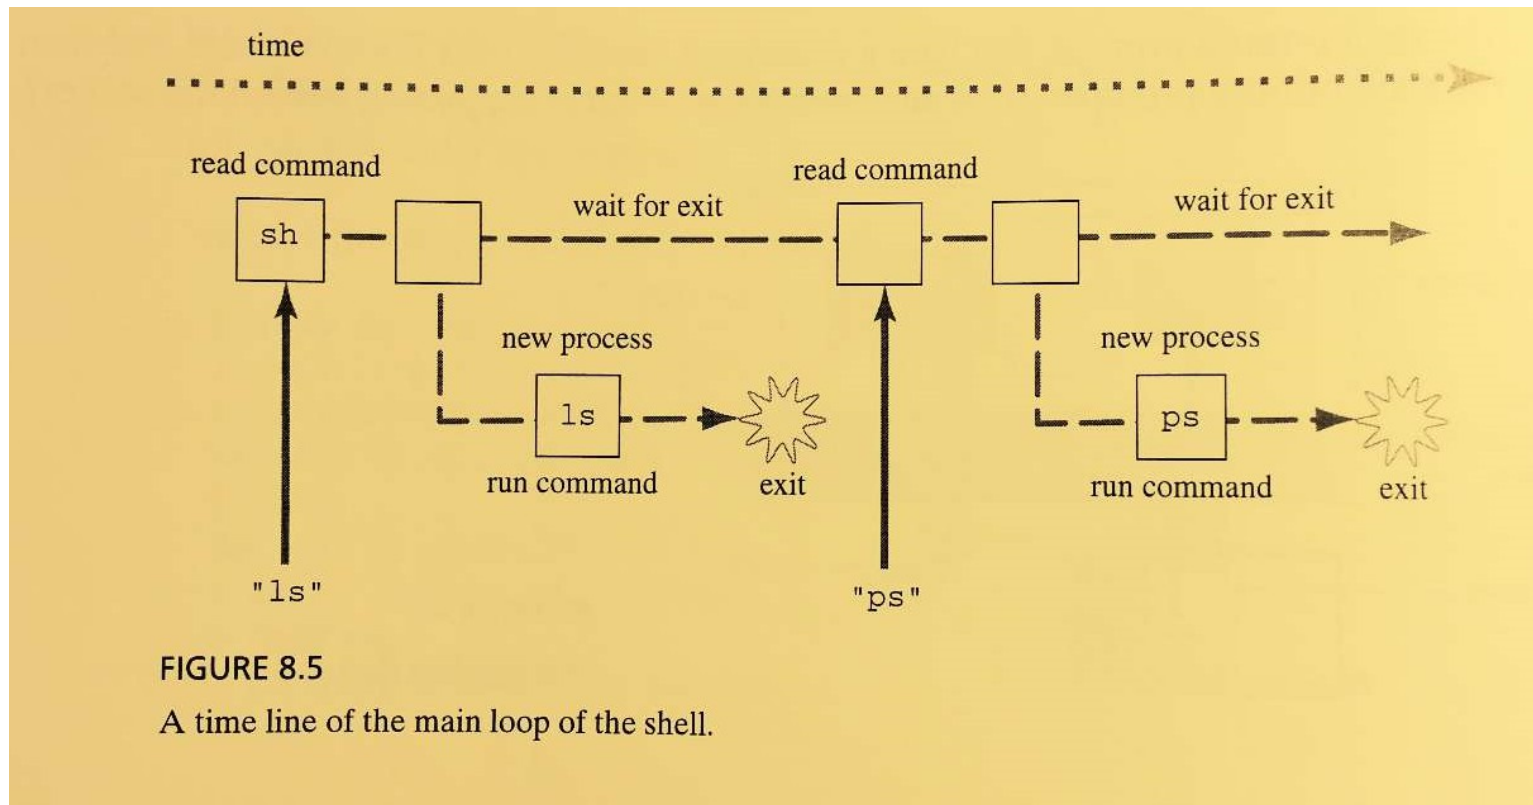
The Main Loop of a Shell

- Consider this typical shell interaction:

```
hstalica@cat:~$ ls
Desktop      etc.listing  Pictures    Public
Documents    examples.desktop  print.c     Templates
Downloads    Music        print.c~    Videos
hstalica@cat:~$ ps
  PID TTY          TIME CMD
 3037 pts/2        00:00:00 bash
 3220 pts/2        00:00:00 ps
hstalica@cat:~$
```

The Main Loop of a Shell

- Let's represent the events with the time line in fig 8.5:



The Main Loop of a Shell

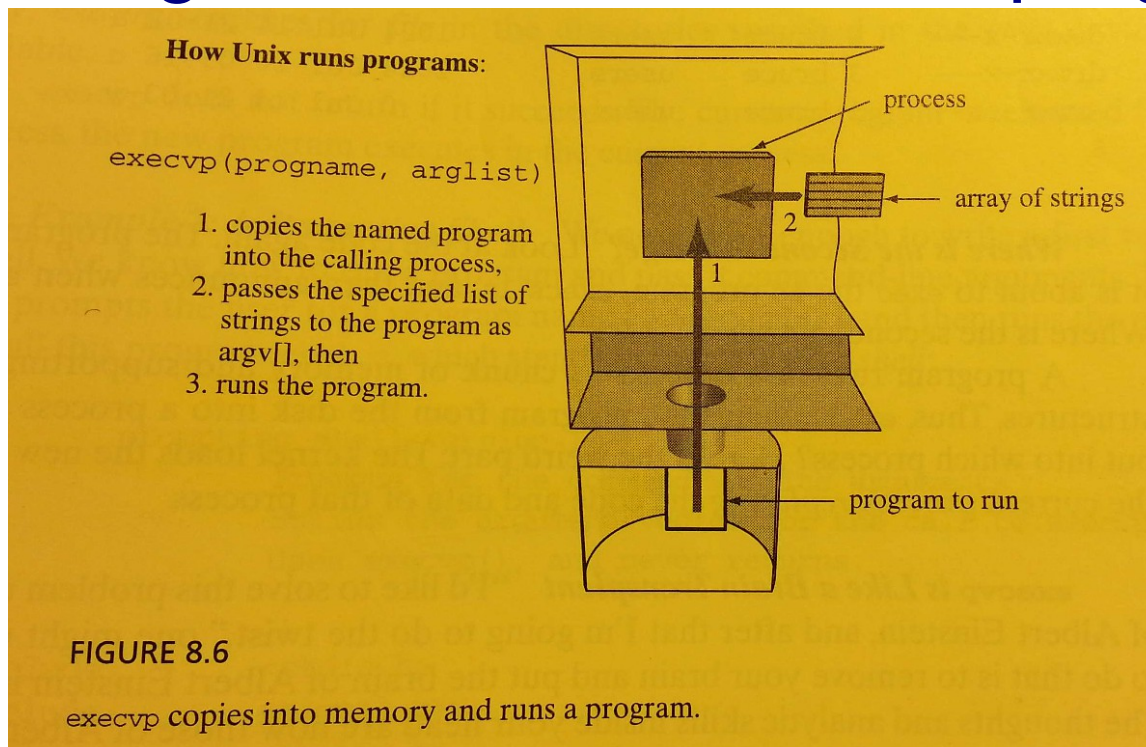
- Time's moving left to right. The shell is the box sh.
- Beginning on the left side, the shell moves right as time passes.
- The shell reads "ls" from the user, creates a new process, then runs the ls command in the process, and waits for it to finish.

The Main Loop of a Shell

- The shell then reads new input, creates a new process, runs the program in that process, then waits for that process to finish.
- When the shell detects end of input, it exits.
- To write a shell, we need to know how to:
 1. Run a program.
 2. Create a process.
 3. Wait for `exit()`

Q1: How Does a Program Run a Program?

- Answer: call `execvp`.
- Figure 8.6 shows how a program runs another:



Q1: How Does a Program Run a Program?

- To run *ls -la*, a program calls *execvp("ls", arglist)* where *arglist* is an array of command-line strings.
- The kernel loads the program from disk into memory.
- The command-line arguments *ls* and *-la* are passed to the program, and it runs.

Q1: How Does a Program Run a Program?

- In summary,
 1. Program calls *execvp*.
 2. Kernel loads program into the process.
 3. Kernel copies *arglist* into the process.
 4. Kernel calls *main(argc, argv)*

Q1: How Does a Program Run a Program?

- Let's look at a program that runs `ls -l`, `exec1.c`.
- Notice `execvp` takes two arguments: the name of the program to run, and the array of command-line arguments for that program.
- The array appears as `argv[]` when the program runs.

Q1: How Does a Program Run a Program?

- We set the first string to the program name.
- Also notice the array must have a null pointer in the last element.
- Compile and run:

```
hstalica@cat:~/Desktop/CS3560/CH08$ cc exec1.c -o exec1
hstalica@cat:~/Desktop/CS3560/CH08$ ./exec1
* * * About to exec ls -l
total 16
-rwxrwxr-x 1 hstalica hstalica 8632 May 17 10:05 exec1
-rw-rw-rw- 1 hstalica hstalica  426 May 17 10:04 exec1.c
hstalica@cat:~/Desktop/CS3560/CH08$
```


Where's the second message?

- Review the code. The program announced it will *exec /s*, then it *execs /s*, and then announces when *exec* finishes. But where's the second message?
- A program runs in a process, so *execvp* loads the program into a process, but which process?

Where's the second message?

- Well, the kernel loads the new program into the current process, *replacing* the code and data of that process!

execvp is Like a Brain Transplant

- The `exec` system call clears out the code of the current program from the current process, then, in the now empty process, puts the code of the program named in the `exec` call, then runs the new program.
- `exec` changes the memory allocation of the process to fit the new program requirements.
- The process is the same, the content is new.

Summary of execvp()

PURPOSE Execute a file, with PATH searching

INCLUDE #include <unistd.h>

USAGE result = execvp(const char* file, const char* args[])

ARGS file name of file to execute

args array of C-Strings

RETURNS -1 if error

Summary of `execvp()`

- `execvp` loads program specified by *file* into the current process and attempts to execute.
- `execvp` passes a list of strings in a NULL-terminated array *argv* to the program.
- `execvp` searches for *file* in directories specified in `PATH`.
- `execvp` doesn't return if successful. The current program is removed, the new program executes in the current process.

Q2: How Do We Get a New Process?


- Answer: Have the process call *fork* to replicate itself.
- Usage: *fork()*; // takes no arguments

fork Explanation


- The solution is to clone the process that calls *execvp* and then have the clone call *execvp*, leaving the original process alone.
- This way, it's the clone shell that dies, and the original shell lives on.

fork Explanation

Before fork:



After fork:



The new process contains the same code and data as the parent process.

FIGURE 8.8

`fork()` makes a copy of a process.

fork Explanation

- Fig 8.8 shows the system before and after the *fork* call.
- The process contains the program and a current location in the program.
- The process calls *fork*.
- Control passes into the *fork* code inside the kernel.

fork Explanation

- The kernel does the following:
 - (a) allocates a new chunk of memory and kernel data structures
 - (b) copies original process into the new one
 - (c) adds the new process to the set of running processes
 - (d) returns control back to *both* processes

fork Explanation

- After cloning the process, there are two processes, both identical, both in the middle of the same instruction.
- However, each is a separate, new process, each able to continue execution independently.
- Let's consider *forkdemo1.c*

fork Explanation

- Were this a normal program, we'd see two lines of output, one for each print statement. However, when run we see


```
hstalica@cat:~/Desktop/CS3560/CH08$ cc forkdemo1.c -o forkdemo1
hstalica@cat:~/Desktop/CS3560/CH08$ ./forkdemo1
Before: my pid is 3355
After: my pid is 3355, fork() said 3356
After: my pid is 3356, fork() said 0
hstalica@cat:~/Desktop/CS3560/CH08$
```

fork Explanation

- We see 3 lines of output, one Before: message and two After: messages.
- Process 3335 prints a Before: and it prints an After:.
- Process 3356 prints the other After: message, but not Before: message. Why?


fork Explanation

Before fork:



One flow of control enters the fork kernel code.

After fork:



Two flows of control return from fork kernel code.

FIGURE 8.9

The child executes the code after `fork()`.

fork Explanation

- The kernel creates process 3356 by cloning process 3355, copying the code and the *current line* in the code into the new process.
- The new process, 3356, picks up in the middle, and does not print a Before: message.

Example: forkdemo2.c – Children Creating Processes

- The child process begins its life, not at the start of *main* but at the return from fork.
- Can you predict how many lines of output are produced?

Example: forkdemo3.c – Distinguishing Parent from Child

- In *forkdemo1.c*, process 3355 called *fork*, creating child process 3356.
- Both processes ran the same code at the same line with all the same data and process attributes.
- How can a process tell if it's the parent or the child?

Example: forkdemo3.c – Distinguishing Parent from Child

- The two processes are not identical.
- Output from *forkdemo1.c* shows *fork* returns different values to the different processes.
- *fork* returns 0 in the child process.
- *fork* returns the PID of the child process in the parent process (3356 in this case).

Example: forkdemo3.c – Distinguishing Parent from Child

- Checking *fork*'s return value is the easiest way to figure out if you are the child or the parent.
- *forkdemo3.c* shows how a single program uses *fork*'s return value to print different messages.

Example: forkdemo3.c – Distinguishing Parent from Child

- A sample run:

```
hstalica@cat:~/Desktop/CS3560/CH08$ cc forkdemo3.c -o forkdemo3
hstalica@cat:~/Desktop/CS3560/CH08$ ./forkdemo3
Before: my pid is 3372
I am the parent, my child is 3373
hstalica@cat:~/Desktop/CS3560/CH08$ I am the child. my pid=3373
```

Summary of *fork*

PURPOSE : Create a process

INCLUDE : `#include <unistd.h>`

USAGE : `pid_t result = fork(void)`

ARGS : None

RETURNS :

-1	if error
0	to child process
pid	pid of child to parent process

Summary of *fork*

- Using *fork*, we create a new process, and can tell which is the new and which is the original process.
- The new process then calls *execvp*.

Q3: How Does the Parent Wait for the Child to Exit?

- Answer: A process calls *wait* to wait for a child to finish.
- Usage: *pid = wait(&status);*

wait Explanation

- *wait* does two things:
 1. it pauses the calling process until the child process finishes.
 2. it retrieves the value the child process passed to *exit*.

wait Explanation

- Fig 8.10 shows us how it works:

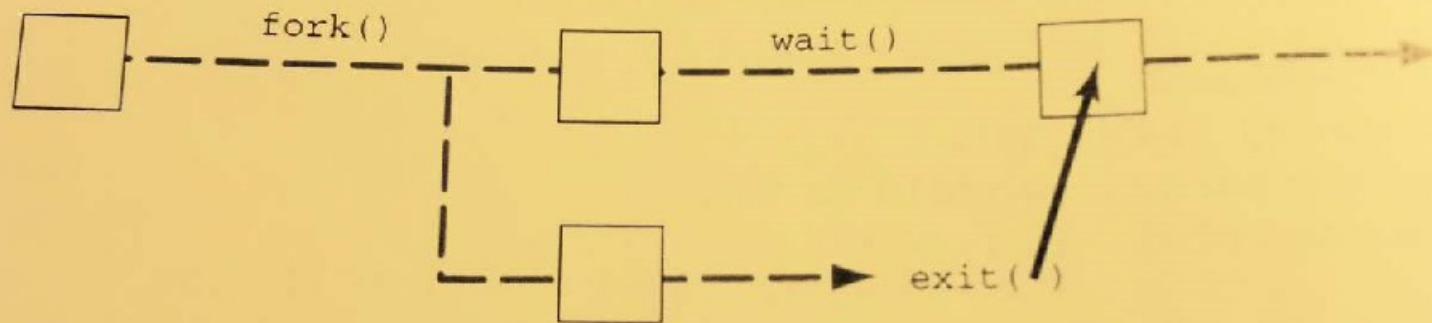


FIGURE 8.10

`wait` pauses the parent until the child finishes.

wait Explanation

- Time flows left to right.
- The parent begins at the left and calls *fork*.
- The kernel constructs the child process and starts running it concurrently with its parent.
- The parent calls *wait* and the kernel suspends the parent until the child finishes
- The parent pauses for the part of the diagram marked *wait*.

wait Explanation

- Eventually, the child process finishes and calls *exit(n)*, passing a number between 0 and 255 as an argument.
- When the child calls *exit*, the kernel wakes up the parent and delivers the argument.
- This notification and transfer is represented by arrow leading from the *exit* parentheses to the parent process.

wait Explanation

- So, *wait* does two things: notification and communication.

Example: waitdemo1.c – notification

- *waitdemo1.c* shows how the *exit* call in the child process triggers a return from *wait* in its parent.
- A sample run:

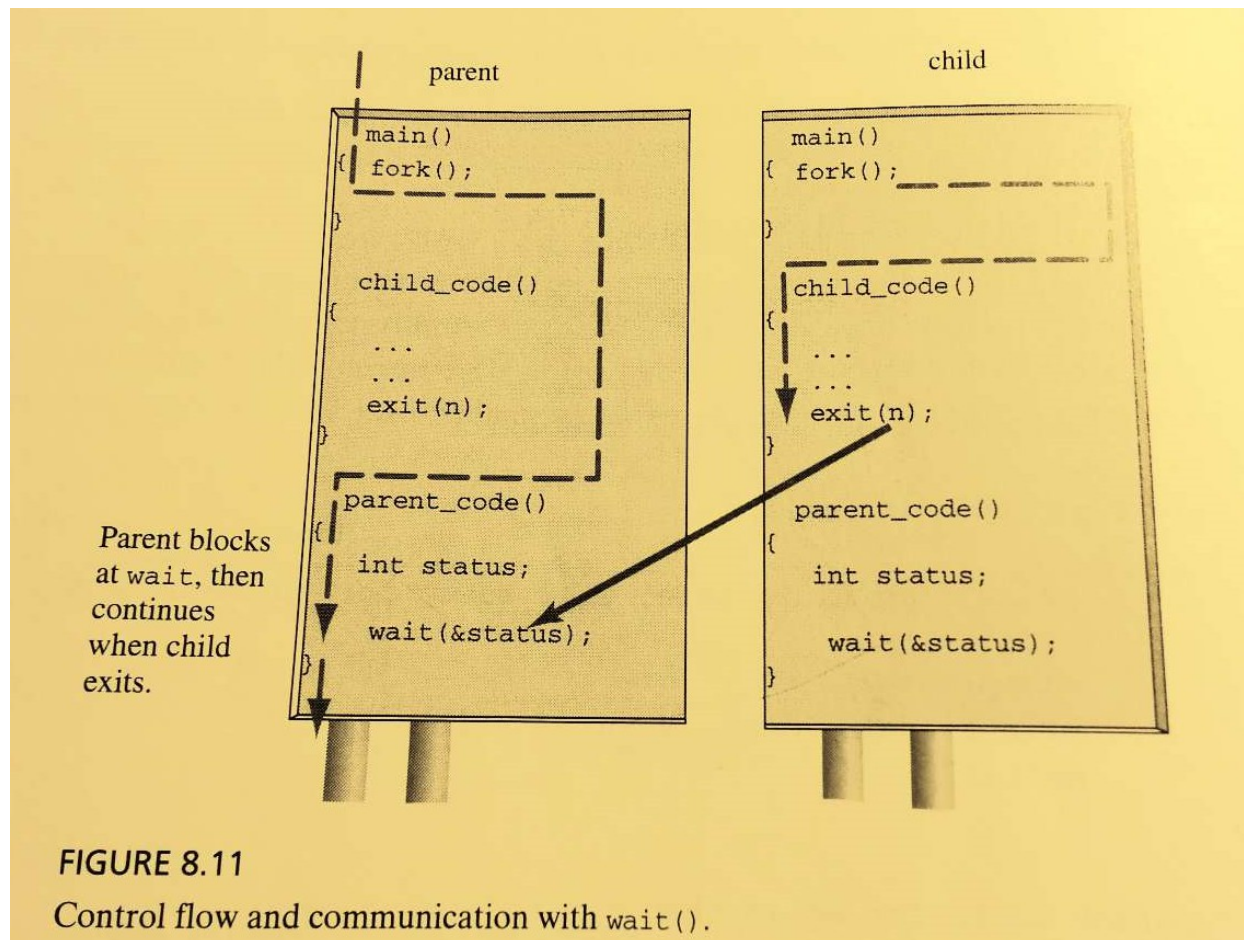
```
hstalica@cat:~/Desktop/CS3560/CH08$ cc waitdemo1.c -owaitdemo1
hstalica@cat:~/Desktop/CS3560/CH08$ ./waitdemo1
before: mypid is 3416
child 3417 here. will sleep for 2 seconds
done waiting for 3417. Wait returned: 3417
4352
hstalica@cat:~/Desktop/CS3560/CH08$
```

Example: waitdemo1.c – notification

- Run the program and adjusting the time will show the parent always waits until the child calls *exit*.
- Fig 8.11 shows flow of control and data transfer between the two.
- In the parent, flow control starts in the middle of *main*, continues to the *child_code* function, and ends with the call to *exit*.

Example: waitdemo1.c – notification

- The *exit* call is like a wake up signal to the parent.



Conclusions from *waitdemo1.c*

- *wait* blocks the calling program until a child finishes.

A parent could *fork* a child process to mow the lawn, then wait until it's done.

The pair of calls *exit* and *wait* is a way to synchronize these tasks.

- *wait* returns the *PID* of the finishing process. Return value from *wait* can tell which of several tasks is finished, so it can continue.

Example: *waitdemo2.c* - communication

- One use of *wait* is to notify the parent that a child is done. The other is to tell the parent how the child finished.
- A process ends in one of three ways: success (0), failure (nonzero value), death (by signal).
- *wait* returns to the parent the PID of the child process that finished. How does the parent know what happened?

Example: *waitdemo2.c* - communication

- Answer: the argument to *wait*.
- The parent calls *wait* with an integer variable address, the kernel stores in that integer the termination status of the child.
- If the child calls *exit* the kernel puts the exit value into the integer.
- If the child is killed, the kernel puts the signal number into the integer.

Example: *waitdemo2.c* - communication

- The integer consists of 3 regions: 8 bits for exit value, 7 bits for the signal number, and 1 bit to indicate a core dump.

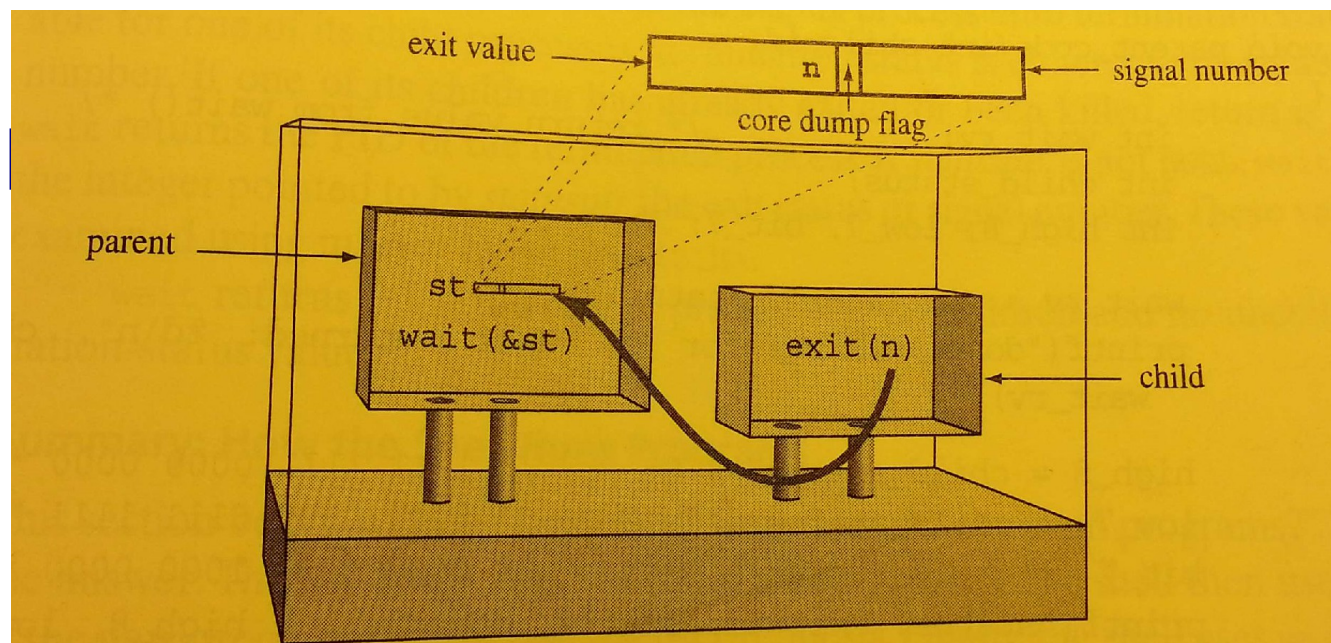


FIGURE 8.12

The child status value has three parts.

Example: *waitdemo2.c* - communication

- Let's look at *waitdemo2.c*.
- First, let's allow it to exit normally:

```
hstalica@cat:~/Desktop/CS3560/CH08$ cc waitdemo2.c -owaitdemo2
hstalica@cat:~/Desktop/CS3560/CH08$ ./waitdemo2
before: mypid is 3439
child 3440 here. will sleep for 5 seconds
child done, about to exit
done waiting for 3440. Wait returned: 3440
status: exit=17, sig=0, core=0
hstalica@cat:~/Desktop/CS3560/CH08$
```

Example: *waitdemo2.c* - communication

- Then, run it in the background and then use the kill command to send SIGTERM to the child:

```
hstalica@cat:~/Desktop/CS3560/CH08$ ./waitdemo2 &  
[1] 3453  
hstalica@cat:~/Desktop/CS3560/CH08$ before: mypid is 3453  
child 3454 here. will sleep for 5 seconds  
kill 3454  
hstalica@cat:~/Desktop/CS3560/CH08$ done waiting for 3454. Wait re  
turned: 3454  
status: exit=0, sig=15, core=0
```

Summary of *wait()*

PURPOSE :	wait for process termination	
INCLUDE :	#include<sys/types.h> #include<sys/wait.h>	
USAGE :	pid_t result = wait(int* statusptr)	
ARGS :	statusptr child result	
RETURNS :	-1	if error
	pid	of terminated process
SEE ALSO :	waitpid (2), wait3 (2)	

Summary of *wait()*

- *wait()* suspends the calling process until termination status is available for one of its child processes.
- Termination status is either an exit value or a signal number.
- If a child has exited or been killed already, return is immediate.

Summary of *wait()*

- *wait* returns the PID of the terminated process.
- If *statusptr* is not NULL, *wait* copies into the integer pointed to by *statusptr* the exit status or signal number.
- These values can be examined using macros in `<sys/wait.h>`
- *wait* returns -1 if the calling process has no children and no uncollected termination-status values.

Summary: How the Shell Runs Programs

- The shell uses *fork* to create a new process.
- The shell then uses *exec* to run, in the new process, the program the user requests.
- Finally, the shell uses *wait* to wait until the process finishes running the command.
- *wait* also obtains the exit status or signal number telling how the child died from the kernel.

Summary: How the Shell Runs Programs

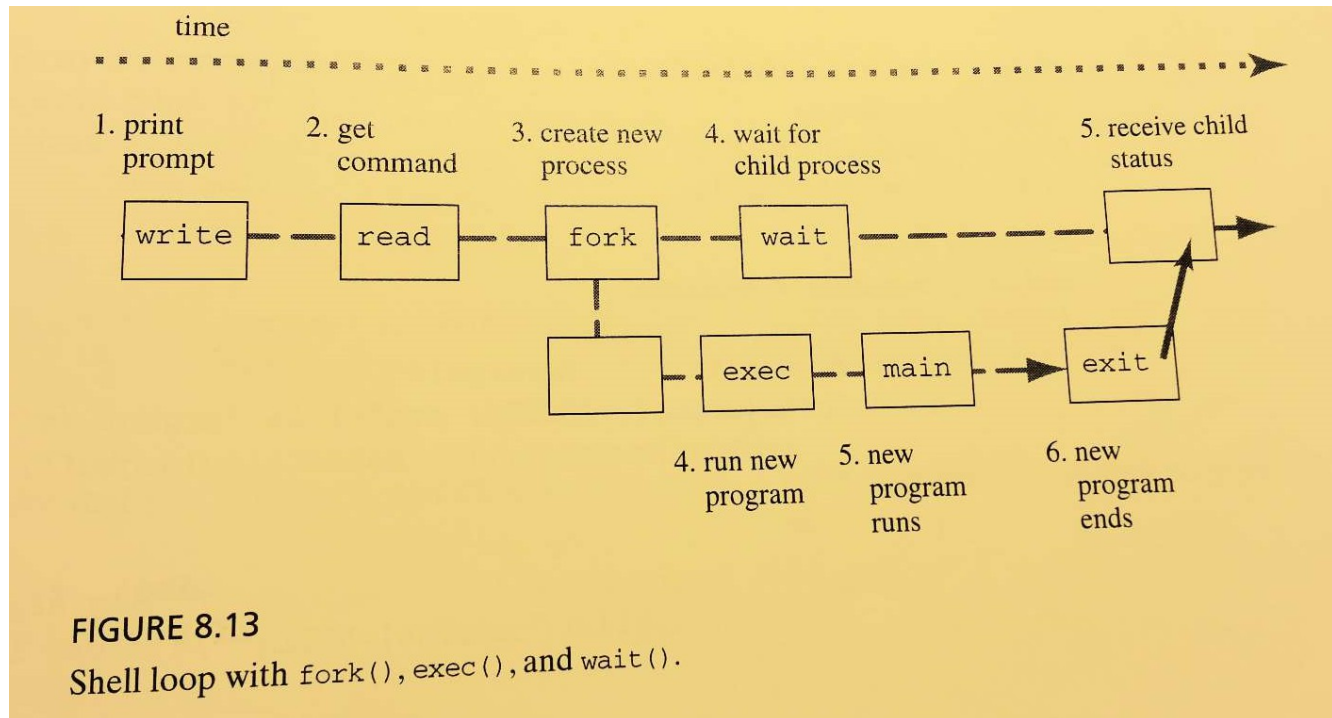


FIGURE 8.13
Shell loop with `fork()`, `exec()`, and `wait()`.

- Every Unix shell uses this model. Let's combine the 3 calls to write a real shell.

Summary

Summary

- Unix runs a program by loading the executable code into a process and running that code. A process is memory space and other system resources required to run a program.
- Each running program runs in its own process. A process has a unique process ID number, and owner, a size, and other properties.

Summary

- The *fork* system call creates a new process by making an almost exact replica of the calling process. The new process is called a child process.
- A program loads and executes a new program in the current process by calling a function in the *exec* family. We studied *execvp*, but there are others.

Summary

- A program can wait for a child process to terminate by calling *wait*.
- A calling program can pass a list of strings to *main* in the new program. The new program can send back a small integer value by calling *exit*.
- A Unix shell runs programs by calling *fork*, *execvp*, and *wait*.

Free eye bleach

